

Intermediate tidyverse exercises

RUG @ HSG

Below, you will find some exercises to apply and improve your tidyverse skills. Generally, they increase in difficulty as you go along. Don't hesitate to ask questions in the Q&A WhatsApp group, we or your fellow students will be happy to answer them!

The data set

In these exercises we will be working with the `gapminder` data set. It includes historical data on the life expectancy, population, and GDP in different countries over the past decades. You can load it directly from the `gapminder` package like this:

```
install.packages("gapminder")  
gapminder <- gapminder::gapminder
```

Task 1: Getting an overview over the data

- (a) In `tidyverse` there are multiple functions that allow you to gain a quick overview over the data that you are working with. Use one of them to find out:
 - How many rows and columns are in the data
 - Which variable types are in the data
- (b) Find out for which time frame the data is available.
- (c) Count the number of countries represented per continent. (*Hint*: the function `n_distinct()` gives you the number of unique values in a vector)

Task 2: Descriptive statistics

- (a) We have the GDP per capita for a lot of countries, but how do they compare per continent? Calculate the mean GDP per capita for every continent in the year 2007. The mean GDP can be misleading if the distribution over the countries is skewed. Also add the median GDP per capita.
- (b) The economic development of these continents seems to be quite different. Calculate the minimum, median, and maximum life expectancy per continent in 1952. Sort them from lowest to highest median life expectancy.
- (c) The table from task (b) is great but a bit hard to read at a glance. Calculate the distance from the minimum life expectancy to the median and from the median to the maximum in the year 1952.



Task 3: More in-depth analysis

- (a) So far we have compared certain variables in one year, now we would like to know how they developed over time. Calculate the median GDP per capita for every continent for every year since 1980. To make the resulting table easier to interpret, transform it to wide format.

Hint: Instead of typing out all years since 1980, try to use a suitable logical condition when filtering.

- (b) We have looked at data on the continent level but how do things look in Switzerland (or your home country)? Extract the GDP per capita, population and life expectancy in Switzerland for all years since 1980. To make the table more readable, display the GDP per capita in thousands, the population in millions, and the life expectancy in years past 50. Transform your result into wide format, so that there is a row per calculated value and a column per year.

Hint: You will have to transform to a "real" long format first, before transforming to wide format.

- (c) Now that we have analyzed Switzerland (or your home country), let's compare it to the rest of Europe (or your home continent). Assemble a table with one column for Switzerland's GDP per capita for every year and one column with Europe's (excluding Switzerland) median GDP per capita for every year.

Hint: You can pass entire "pipe chains" as arguments to functions in other pipe chains.

If you made it this far, we commend you! That was some serious data wrangling just then. You now have some important skills to apply to your own data analyses!

Solutions

Task 1

- (a) There is more than one answer to this task, however, we prefer `glimpse()` because the overview is nice and condensed.

```
gapminder %>% glimpse()
```

```
## Rows: 1,704
## Columns: 6
## $ country   <fct> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", ~
## $ continent <fct> Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, ~
## $ year      <int> 1952, 1957, 1962, 1967, 1972, 1977, 1982, 1987, 1992, 1997, ~
## $ lifeExp   <dbl> 28.801, 30.332, 31.997, 34.020, 36.088, 38.438, 39.854, 40.8~
## $ pop       <int> 8425333, 9240934, 10267083, 11537966, 13079460, 14880372, 12~
## $ gdpPercap <dbl> 779.4453, 820.8530, 853.1007, 836.1971, 739.9811, 786.1134, ~
```

We see that there are 1704 rows and 6 columns in the dataframe. We have factor, integer and double data types.

- (b) We can use the `summarize()` function for this:

```
gapminder %>%
  summarize(starting_year = min(year),
            ending_year = max(year))
```

```
## # A tibble: 1 x 2
##   starting_year ending_year
##   <int>         <int>
## 1      1952         2007
```

- (c) If we group by continent first, we can use the `n_distinct()` function to give us the number of countries per continent.

```
gapminder %>%
  group_by(continent) %>%
  summarize(nr_countries = n_distinct(country))
```

```
## # A tibble: 5 x 2
##   continent nr_countries
##   <fct>         <int>
## 1 Africa             52
## 2 Americas           25
## 3 Asia               33
## 4 Europe             30
## 5 Oceania            2
```



Task 2

- (a) To analyze only the year 2007 and not the entire time frame, we first filter for the year 2007 before we group the countries.

```
gapminder %>%  
  filter(year == 2007) %>%  
  group_by(continent) %>%  
  summarize(mean_gdpPercap = mean(gdpPercap),  
            median_gdpPercap = median(gdpPercap))
```

```
## # A tibble: 5 x 3  
##   continent mean_gdpPercap median_gdpPercap  
##   <fct>          <dbl>          <dbl>  
## 1 Africa          3089.          1452.  
## 2 Americas       11003.          8948.  
## 3 Asia           12473.          4471.  
## 4 Europe         25054.         28054.  
## 5 Oceania        29810.         29810.
```

- (b) Similarly to the task before, we can calculate these values by filtering, grouping, and summarizing.

```
gapminder %>%  
  filter(year == 1952) %>%  
  group_by(continent) %>%  
  summarize(min_lifeExp = min(lifeExp),  
            median_lifeExp = median(lifeExp),  
            max_lifeExp = max(lifeExp)) %>%  
  arrange(median_lifeExp)
```

```
## # A tibble: 5 x 4  
##   continent min_lifeExp median_lifeExp max_lifeExp  
##   <fct>          <dbl>          <dbl>          <dbl>  
## 1 Africa          30          38.8          52.7  
## 2 Asia           28.8         44.9          65.4  
## 3 Americas       37.6         54.7          68.8  
## 4 Europe         43.6         65.9          72.7  
## 5 Oceania        69.1         69.3          69.4
```



(c) We can pass more complex calculations with more than one function to `summarize`.

```
gapminder %>%
  filter(year == 1952) %>%
  group_by(continent) %>%
  summarize(neg_dev_lifeExp = median(lifeExp) - min(lifeExp),
            pos_dev_lifeExp = max(lifeExp) - median(lifeExp))
```

```
## # A tibble: 5 x 3
##   continent neg_dev_lifeExp pos_dev_lifeExp
##   <fct>          <dbl>          <dbl>
## 1 Africa           8.83           13.9
## 2 Americas        17.2           14.0
## 3 Asia            16.1           20.5
## 4 Europe          22.3            6.77
## 5 Oceania          0.135          0.135
```

Note: The deviation in Oceania is so small because there are only two countries included from this continent (Australia and New Zealand). We know this from our calculation in task 1 (c).

Task 3

(a) You already know how to calculate the medians from previous tasks. Additionally, we can transform the data to wide format with `pivot_wider()`.

```
gapminder %>%
  filter(year >= 1980) %>%
  group_by(continent, year) %>%
  summarize(median_gdpPercap = median(gdpPercap)) %>%
  pivot_wider(names_from = year,
              values_from = median_gdpPercap)
```

```
## `summarise()` has grouped output by 'continent'. You can override using the
## `.groups` argument.
```

```
## # A tibble: 5 x 7
## # Groups:   continent [5]
##   continent `1982` `1987` `1992` `1997` `2002` `2007`
##   <fct>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Africa    1324. 1220. 1162. 1180. 1216. 1452.
## 2 Americas  6435. 6361. 6619. 7114. 6995. 8948.
## 3 Asia      4107. 4106. 3726. 3645. 4091. 4471.
## 4 Europe   15323. 16215. 17550. 19596. 23675. 28054.
## 5 Oceania   18555. 20448. 20894. 24024. 26939. 29810.
```

Note: The logical operator `>=` is equivalent to the \geq sign. Similarly, we can use `<=` to signify \leq . We can also use the logical "not-equal-to" operator `!=` to mean \neq .



- (b) To preserve the year variable, we group by it first, even though there is only one value in each of the groups. We could also use `mutate()` to add the values back in but then we would have to type out all the different years (way too annoying!),

```
gapminder %>%
  filter(country == "Switzerland" & year >= 1980) %>%
  group_by(year) %>%
  summarize(k_gdpPercap = gdpPercap / 1e3,
            mio_pop = pop / 1e6,
            lifeExp_past_50 = lifeExp - 50) %>%
  ungroup() %>%
  pivot_longer(cols = c(k_gdpPercap, mio_pop, lifeExp_past_50),
               names_to = "variables",
               values_to = "values") %>%
  pivot_wider(names_from = year,
              values_from = values)
```

```
## # A tibble: 3 x 7
##   variables    `1982` `1987` `1992` `1997` `2002` `2007`
##   <chr>        <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 k_gdpPercap  28.4  30.3  31.9  32.1  34.5  37.5
## 2 mio_pop      6.47  6.65  7.00  7.19  7.36  7.55
## 3 lifeExp_past_50 26.2  27.4  28.0  29.4  30.6  31.7
```

Note: The `1e3` and `1e6` are scientific notation for 1'000 and 1'000'000. By using this, you can avoid mistakes coming from typing too many or too few zeros.

- (c) The key to solving this task in one “pipe chain” is that you can pass pipe chains as arguments within other pipe chains as if they were just variables. If you solved this by first saving the Europe column in a separate variable, this is of course completely valid. This applies to coding in general: everyone has their own preferred way to do things and everyone will defend that approach as if their life depended on it.

```
gapminder %>%
  filter(country == "Switzerland") %>%
  group_by(year) %>%
  summarize(k_gdpPercap_CH = gdpPercap / 1e3) %>%
  mutate(gapminder %>%
    filter(continent == "Europe" & country != "Switzerland") %>%
    group_by(year) %>%
    summarize(k_median_gdpPercap_EU = median(gdpPercap) / 1e3) %>%
    select(k_median_gdpPercap_EU)
  )
```

```
## # A tibble: 12 x 3
##   year k_gdpPercap_CH k_median_gdpPercap_EU
##   <int>      <dbl>          <dbl>
## 1  1952         14.7           5.07
## 2  1957         17.9           6.04
## 3  1962         20.4           7.48
## 4  1967         23.0           9.33
## 5  1972         27.2          12.3
## 6  1977         27.0          14.2
## 7  1982         28.4          15.3
## 8  1987         30.3          16.1
## 9  1992         31.9          17.5
## 10 1997         32.1          18.7
## 11 2002         34.5          22.5
## 12 2007         37.5          27.5
```